

Background

Learning when to use a decomposition,
Markus Kruber et al.

motivation

algorithm

original MIP


$\min \{c^t x, Ax \geq b, x \in \mathbb{Z}_+^n \times \mathbb{Q}_+^q\}$. It can sometimes be re-arranged such that a particular structure:

$$\begin{array}{ll} \min & c^t x \\ \text{s.t.} & \begin{bmatrix} D^1 & & & F^1 \\ & D^2 & & F^2 \\ & & \ddots & \vdots \\ & & & D^\kappa & F^\kappa \\ A^1 & A^2 & \dots & A^\kappa & G \end{bmatrix} \cdot \begin{bmatrix} x^1 \\ x^2 \\ \vdots \\ x^\kappa \\ x^\ell \end{bmatrix} \geq \begin{bmatrix} b^1 \\ b^2 \\ \vdots \\ b^\kappa \\ b^\ell \end{bmatrix} \\ & x \in \mathbb{Z}_+^n \times \mathbb{Q}_+^q. \end{array}$$

Such a re-arrangement is called a **decomposition** of the original MIP and finding it is called **detection**.

The MIP can be reformulated according to “best suited” decomposition, if the MIP structure can be automatically and finely detected. As a result, the solver may be much faster on the reformulated model than on the original one.

Dantzig-Wolfe (DW) reformulation is very efficient on solving specially structured MIPs. Successfully applying it may require a solid background, experience, and a non-negligible implementation effort. This paper proposes a supervised learning approach to decide whether or not a reformulation should be applied, and which decomposition to choose when several are possible. Preliminary experiments with a MIP solver equipped with this knowledge show a significant performance improvement on structured instances, with little deterioration on others.



SCIP is a well-established MIP solver, **GCG** is its extension to make DW reformulation more accessible.

Their Supervised Learning Approach: They would like to learn an answer to the question: Given a MIP P , a DW decomposition D , and a time limit τ , will GCG using D optimally solve P faster than SCIP? They define a mapping ϕ that transforms a tuple $(P; D; \tau)$ into a vector of sufficient statistics or features $\phi(P; D; \tau)$. Due to this mapping ϕ , the question above becomes a standard binary classification problem. Therefore a standard classifier $f: \{0,1\}$ can be trained to solve this problem. Given an instance $(P; D; \tau)$, the quantity $f \circ \phi(P; D; \tau)$ is equal to one iff the predicted answer to the question above is positive. Practically, they built a database of SCIP and GCG runs for tuples $(P; D; \tau)$, a mapping ϕ , and trained classifiers f from the scikit-learn library on the instances $\phi(P; D; \tau)$. Answers to the probabilistic versions $g: [0, 1]$ of these classifiers can be interpreted as the probability that GCG using D outperforms SCIP if the time limit is τ .